

# Loss Functions and Numerical Optimization

Ravi Kothari, Ph.D.  
ravi.kothari@ashoka.edu.in

“If you cannot measure it, you cannot improve it - Lord Kelvin”

## Learning – A Formal Definition

*Based on  $N$  (possibly noisy) observations  $\mathcal{X} = \{(x^{(i)}, y^{(i)})\}_{i=1}^N$  of the input and output of a fixed though unknown system  $f(x)$ , construct an estimator  $\hat{f}(x; \theta)$  so as to minimize,*

$$E \left[ \left( L(f(x) - \hat{f}(x; \theta)) \right) \right]$$

- $L$  is the *loss function* and specifies *how* we want to measure the difference between the “desired output” ( $f(x)$ ) and the “actual output” ( $\hat{f}(x; \theta)$ )

## Learning – A Formal Definition

*Based on  $N$  (possibly noisy) observations  $\mathcal{X} = \{(x^{(i)}, y^{(i)})\}_{i=1}^N$  of the input and output of a fixed though unknown system  $f(x)$ , construct an estimator  $\hat{f}(x; \theta)$  so as to minimize,*

$$E \left[ \left( L(f(x) - \hat{f}(x; \theta)) \right) \right]$$

- $L$  is the *loss function* and specifies *how* we want to measure the difference between the “desired output” ( $f(x)$ ) and the “actual output” ( $\hat{f}(x; \theta)$ )
- We have seen a loss function already; recall our discussion when we did the Bayes classifier

# Some Loss Functions

## Some Loss Functions

- *Classification*:  $f(x) \in \{\omega_1, \omega_2, \dots, \omega_c\}$

$$L(f(x), \hat{f}(x; \theta)) = \begin{cases} 0, & f(x) = \hat{f}(x; \theta) \\ 1, & f(x) \neq \hat{f}(x; \theta) \end{cases}$$

# Some Loss Functions

- *Classification*:  $f(x) \in \{\omega_1, \omega_2, \dots, \omega_c\}$

$$L(f(x), \hat{f}(x; \theta)) = \begin{cases} 0, & f(x) = \hat{f}(x; \theta) \\ 1, & f(x) \neq \hat{f}(x; \theta) \end{cases}$$

- *Regression*:  $f(x) \in \mathcal{R}^m$

$$L(f(x), \hat{f}(x; \theta)) = \left( f(x) - \hat{f}(x; \theta) \right)^2$$

## Some Loss Functions

- *Classification*:  $f(x) \in \{\omega_1, \omega_2, \dots, \omega_c\}$

$$L(f(x), \hat{f}(x; \theta)) = \begin{cases} 0, & f(x) = \hat{f}(x; \theta) \\ 1, & f(x) \neq \hat{f}(x; \theta) \end{cases}$$

- *Regression*:  $f(x) \in \mathcal{R}^m$

$$L(f(x), \hat{f}(x; \theta)) = \left( f(x) - \hat{f}(x; \theta) \right)^2$$

- *Density Estimation*:

$$L(\hat{f}(x; \theta)) = -\log \hat{f}(x; \theta)$$

## Some Loss Functions

- *Classification*:  $f(x) \in \{\omega_1, \omega_2, \dots, \omega_c\}$

$$L(f(x), \hat{f}(x; \theta)) = \begin{cases} 0, & f(x) = \hat{f}(x; \theta) \\ 1, & f(x) \neq \hat{f}(x; \theta) \end{cases}$$

- *Regression*:  $f(x) \in \mathcal{R}^m$

$$L(f(x), \hat{f}(x; \theta)) = \left( f(x) - \hat{f}(x; \theta) \right)^2$$

- *Density Estimation*:

$$L(\hat{f}(x; \theta)) = -\log \hat{f}(x; \theta)$$

The loss function may not be *symmetric* (in fact, it is not in many cases e.g. in medical diagnosis, a false negative is MUCH worse than a false positive)



# Learning

# Learning

- Once a loss function is defined, learning becomes an optimization process that minimizes the loss function by adjusting the parameters ( $\theta$ )

# Learning

- Once a loss function is defined, learning becomes an optimization process that minimizes the loss function by adjusting the parameters ( $\theta$ )
- This is the *prevalent and popular* definition of learning (though completely inappropriate in my view – more on it later)

# Numerical Optimization

# Numerical Optimization

- The general idea is to start with some values of the parameters, present a training pattern, adjust the parameters to reduce the loss function, present the next training pattern and repeating until the loss function is below an acceptable threshold (or stops decreasing)

# Numerical Optimization

- The general idea is to start with some values of the parameters, present a training pattern, adjust the parameters to reduce the loss function, present the next training pattern and repeating until the loss function is below an acceptable threshold (or stops decreasing)
- i.e.

$$\theta' = \theta + \Delta\theta \quad (1)$$

such that,

$$L(t+1) < L(t) \quad (2)$$

# Numerical Optimization

- The general idea is to start with some values of the parameters, present a training pattern, adjust the parameters to reduce the loss function, present the next training pattern and repeating until the loss function is below an acceptable threshold (or stops decreasing)
- i.e.

$$\theta' = \theta + \Delta\theta \quad (1)$$

such that,

$$L(t+1) < L(t) \quad (2)$$

- A popular approach to finding  $\theta$  is based on *gradient descent*

# Gradient Descent



# Gradient Descent

- Using Taylor Series ( $\theta' = \theta + \Delta\theta$ ),

$$L(\theta') = L(\theta) + \nabla L(\theta)_{\theta}^T \Delta\theta + \dots \quad (3)$$

# Gradient Descent

- Using Taylor Series ( $\theta' = \theta + \Delta\theta$ ),

$$L(\theta') = L(\theta) + \nabla L(\theta)_{\theta}^T \Delta\theta + \dots \quad (3)$$

- To lower the LHS,

$$\nabla L(\theta)_{\theta}^T \Delta\theta < 0 \quad (4)$$

# Gradient Descent

- Using Taylor Series ( $\theta' = \theta + \Delta\theta$ ),

$$L(\theta') = L(\theta) + \nabla L(\theta)_\theta^T \Delta\theta + \dots \quad (3)$$

- To lower the LHS,

$$\nabla L(\theta)_\theta^T \Delta\theta < 0 \quad (4)$$

- This happens maximally when,

$$\Delta\theta \propto -\nabla L(\theta)_\theta \quad (5)$$

# Gradient Descent

- Using Taylor Series ( $\theta' = \theta + \Delta\theta$ ),

$$L(\theta') = L(\theta) + \nabla L(\theta)_{\theta}^T \Delta\theta + \dots \quad (3)$$

- To lower the LHS,

$$\nabla L(\theta)_{\theta}^T \Delta\theta < 0 \quad (4)$$

- This happens maximally when,

$$\Delta\theta \propto -\nabla L(\theta)_{\theta} \quad (5)$$

- i.e.,

$$\Delta\theta = -\eta \nabla L(\theta)_{\theta} \quad (6)$$

where,  $\eta$  is the step size. So,

$$\theta' = \theta - \eta \nabla L(\theta)_{\theta} \quad (7)$$

# Example

## Example

- Say we want to find the minimum of,

$$L(\theta) = \theta_1^2 + (\theta_2 - 3)^2 \quad (8)$$

## Example

- Say we want to find the minimum of,

$$L(\theta) = \theta_1^2 + (\theta_2 - 3)^2 \quad (8)$$

- Let the starting guess be,

$$\theta = \begin{bmatrix} \theta_1 \\ \theta_2 \end{bmatrix} = \begin{bmatrix} 1 \\ 0.7 \end{bmatrix} \quad (9)$$

## Example

- Say we want to find the minimum of,

$$L(\theta) = \theta_1^2 + (\theta_2 - 3)^2 \quad (8)$$

- Let the starting guess be,

$$\theta = \begin{bmatrix} \theta_1 \\ \theta_2 \end{bmatrix} = \begin{bmatrix} 1 \\ 0.7 \end{bmatrix} \quad (9)$$

- So,

$$\nabla L(\theta) = \begin{bmatrix} \frac{\partial L(\theta)}{\partial(\theta_1)} \\ \frac{\partial L(\theta)}{\partial(\theta_2)} \end{bmatrix} = \begin{bmatrix} 2\theta_1 \\ 2\theta_2 - 6 \end{bmatrix} \quad (10)$$



## Example

- Say we want to find the minimum of,

$$L(\theta) = \theta_1^2 + (\theta_2 - 3)^2 \quad (8)$$

- Let the starting guess be,

$$\theta = \begin{bmatrix} \theta_1 \\ \theta_2 \end{bmatrix} = \begin{bmatrix} 1 \\ 0.7 \end{bmatrix} \quad (9)$$

- So,

$$\nabla L(\theta) = \begin{bmatrix} \frac{\partial L(\theta)}{\partial \theta_1} \\ \frac{\partial L(\theta)}{\partial \theta_2} \end{bmatrix} = \begin{bmatrix} 2\theta_1 \\ 2\theta_2 - 6 \end{bmatrix} \quad (10)$$

- Using  $\eta = 0.1$  in,

$$\theta(t+1) = \theta(t) - \eta \nabla_{\theta(t)} \quad (11)$$

we get,

$$\theta = \begin{bmatrix} 1 \\ 0.7 \end{bmatrix} = -0.1 \begin{bmatrix} 2\theta_1 \\ 2\theta_2 - 6 \end{bmatrix} = \begin{bmatrix} 1 \\ 0.7 \end{bmatrix} - 0.1 \begin{bmatrix} 2 \\ -4.4 \end{bmatrix} = \begin{bmatrix} 0.8 \\ 1.14 \end{bmatrix} \quad (12)$$

## Example

- Say we want to find the minimum of,

$$L(\theta) = \theta_1^2 + (\theta_2 - 3)^2 \quad (8)$$

- Let the starting guess be,

$$\theta = \begin{bmatrix} \theta_1 \\ \theta_2 \end{bmatrix} = \begin{bmatrix} 1 \\ 0.7 \end{bmatrix} \quad (9)$$

- So,

$$\nabla L(\theta) = \begin{bmatrix} \frac{\partial L(\theta)}{\partial \theta_1} \\ \frac{\partial L(\theta)}{\partial \theta_2} \end{bmatrix} = \begin{bmatrix} 2\theta_1 \\ 2\theta_2 - 6 \end{bmatrix} \quad (10)$$

- Using  $\eta = 0.1$  in,

$$\theta(t+1) = \theta(t) - \eta \nabla_{\theta(t)} \quad (11)$$

we get,

$$\theta = \begin{bmatrix} 1 \\ 0.7 \end{bmatrix} = -0.1 \begin{bmatrix} 2\theta_1 \\ 2\theta_2 - 6 \end{bmatrix} = \begin{bmatrix} 1 \\ 0.7 \end{bmatrix} - 0.1 \begin{bmatrix} 2 \\ -4.4 \end{bmatrix} = \begin{bmatrix} 0.8 \\ 1.14 \end{bmatrix} \quad (12)$$

# Other Optimization Methods

# Other Optimization Methods

- Gradient descent can get trapped in a local minima (like all local search methods)

# Other Optimization Methods

- Gradient descent can get trapped in a local minima (like all local search methods)
- It can be slow to converge. One can use higher order derivatives. They are expensive to compute though provide a better approximation of the error surface and can converge faster

# Other Optimization Methods

- Gradient descent can get trapped in a local minima (like all local search methods)
- It can be slow to converge. One can use higher order derivatives. They are expensive to compute though provide a better approximation of the error surface and can converge faster
- Quasi-Newton's method (such as *Broyden-Fletcher-Goldfarb-Shanno* (BFGS)) do not explicitly compute the *Hessian* but rather approximate it. They usually converge much faster than first-order gradient descent. We will look at some if time permits