

# Generative Adversarial Networks

Ravi Kothari, Ph.D.  
ravi.kothari@ashoka.edu.in

“A wise enemy is better than a foolish friend” – Unknown

# Generative Modeling

# Generative Modeling

- Our goal in **generative modeling** is to learn the distribution from which the true data was sampled

# Generative Modeling

- Our goal in **generative modeling** is to learn the distribution from which the true data was sampled
- We saw an important use-case of generative models when we covered variational auto-encoder *viz.* learning compact representations

# Generative Modeling

- Our goal in **generative modeling** is to learn the distribution from which the true data was sampled
- We saw an important use-case of generative models when we covered variational auto-encoder *viz.* learning compact representations
- Other reasons include,

# Generative Modeling

- Our goal in **generative modeling** is to learn the distribution from which the true data was sampled
- We saw an important use-case of generative models when we covered variational auto-encoder *viz.* learning compact representations
- Other reasons include,
  - ▶ Generate realistic samples e.g. paintings, art – possibly this is one element of creativity

# Generative Modeling

- Our goal in **generative modeling** is to learn the distribution from which the true data was sampled
- We saw an important use-case of generative models when we covered variational auto-encoder *viz.* learning compact representations
- Other reasons include,
  - ▶ Generate realistic samples e.g. paintings, art – possibly this is one element of creativity
  - ▶ We can tell which data is more likely – something that can be very useful in speech recognition

# Generative Modeling

- Our goal in **generative modeling** is to learn the distribution from which the true data was sampled
- We saw an important use-case of generative models when we covered variational auto-encoder *viz.* learning compact representations
- Other reasons include,
  - ▶ Generate realistic samples e.g. paintings, art – possibly this is one element of creativity
  - ▶ We can tell which data is more likely – something that can be very useful in speech recognition
  - ▶ Of course, learning compact representations remains an attractive reason



# Generative Adversarial Networks (GANs)

# Generative Adversarial Networks (GANs)

- GANs are an *implicit generative model* in the sense that a GAN produces samples that confirms to the distribution from which the original data came from. Explicit generative models on the other hand attempt to estimate the actual distribution

# GAN Architecture

# GAN Architecture

- Generate a sample  $z$  from a simple fixed distribution e.g.  $\mathcal{N}(0, 1)$

# GAN Architecture

- Generate a sample  $z$  from a simple fixed distribution e.g.  $\mathcal{N}(0, 1)$
- There is a **generator network** which produces  $x = G(z)$

# GAN Architecture

- Generate a sample  $z$  from a simple fixed distribution e.g.  $\mathcal{N}(0, 1)$
- There is a **generator network** which produces  $x = G(z)$
- In an implicit generative model, the produced samples are indistinguishable from the true (given unlabeled) data

# GAN Architecture

- Generate a sample  $z$  from a simple fixed distribution e.g.  $\mathcal{N}(0, 1)$
- There is a **generator network** which produces  $x = G(z)$
- In an implicit generative model, the produced samples are indistinguishable from the true (given unlabeled) data
- So, GANs also train a **discriminator network** that tries to tell if the data came from the training set (true data) or whether it was generated from the generator network

# GAN Architecture

- Generate a sample  $z$  from a simple fixed distribution e.g.  $\mathcal{N}(0, 1)$
- There is a **generator network** which produces  $x = G(z)$
- In an implicit generative model, the produced samples are indistinguishable from the true (given unlabeled) data
- So, GANs also train a **discriminator network** that tries to tell if the data came from the training set (true data) or whether it was generated from the generator network
- So, the discriminator network produces  $D(x)$  – the probability that  $x$  is real



# GAN Training

# GAN Training

- So, the generator tries to produce something that is realistic enough to fool the discriminator and the discriminator tries to catch the generator. We can thus define a loss function as,

$$J_D = E_{x \sim D}[-\log D(x)] + E_z[-\log(1 - D(G(z)))] \quad (1)$$

# GAN Training

- So, the generator tries to produce something that is realistic enough to fool the discriminator and the discriminator tries to catch the generator. We can thus define a loss function as,

$$J_D = E_{x \sim D}[-\log D(x)] + E_z[-\log(1 - D(G(z)))] \quad (1)$$

- When this **cross-entropy** loss is low, it implies the discriminator can easily tell the real from fake. If the cross-entropy is high, it means the discriminator is not able to tell them apart

# GAN Training

- So, the generator tries to produce something that is realistic enough to fool the discriminator and the discriminator tries to catch the generator. We can thus define a loss function as,

$$J_D = E_{x \sim D}[-\log D(x)] + E_z[-\log(1 - D(G(z)))] \quad (1)$$

- When this **cross-entropy** loss is low, it implies the discriminator can easily tell the real from fake. If the cross-entropy is high, it means the discriminator is not able to tell them apart
- So, for the generator the obvious cost function is to *maximize* the discriminator's cross-entropy, i.e.,

$$J_G = -J_D = \text{const} + E_z[\log(1 - D(G(z)))] \quad (2)$$

Since the generator has no control over the first term, it is just a constant

# GAN Training

# GAN Training

- So, the generator is trying to compute,

$$\arg \max_G \min_D J_D \quad (3)$$

# GAN Training

- So, the generator is trying to compute,

$$\arg \max_G \min_D J_D \quad (3)$$

- This is also referred to as the **minimax formulation** and is an example of a **zero-sum game**

# GAN Training

- So, the generator is trying to compute,

$$\arg \max_G \min_D J_D \quad (3)$$

- This is also referred to as the **minimax formulation** and is an example of a **zero-sum game**
- The discriminator and the generator are trained jointly on their respective cost-functions using back-propagation



# GAN Training

# GAN Training

- There is a difficulty with cross-entropy in this scenario. For example, assume the discriminator is doing well.  $D(G(z))$  is close to 0 in that case and  $J_G$  is also close to 0 even though the scenario implies the generator was doing poorly

# GAN Training

- There is a difficulty with cross-entropy in this scenario. For example, assume the discriminator is doing well.  $D(G(z))$  is close to 0 in that case and  $J_G$  is also close to 0 even though the scenario implies the generator was doing poorly
- We thus often use the modified cost function for the generator given by,

$$J_G = E_z[-\log D(G(z))] \quad (4)$$

# GAN Examples




Figure 1: Our training starts with both the generator (G) and discriminator (D) having a low spatial resolution of  $4 \times 4$  pixels. As the training advances, we incrementally add layers to G and D, thus increasing the spatial resolution of the generated images. All existing layers remain trainable throughout the process. Here  $N \times N$  refers to convolutional layers operating on  $N \times N$  spatial resolution. This allows stable synthesis in high resolutions and also speeds up training considerably. On the right we show six example images generated using progressive growing at  $1024 \times 1024$ .

From Karras et al., 2017

# Demo

`https://poloclub.github.io/ganlab/`